

FoxFS specification

Moritz "LittleFox" Grosch

October 9, 2015

Contents

1	Preface	1
1.1	Why another filesystem?	1
1.2	Design	2
2	Datastructures in FoxFS	2
2.1	superblock	2
2.2	inode	2
2.3	free space chain	4
2.4	Directories	4
2.5	File-metadata	5
3	The layers of FoxFS	5
3.1	Block-layer	5
3.2	Blockchain-layer	5

1 Preface

1.1 Why another filesystem?

This filesystem has to be implemented in JavaScript, so it has to be simple. No journaling, no access control lists, no unix permissions, no owner ... and so on.

Using something like FAT wasn't an option, because of filesize-limits. Just using 64bit attributes makes FAT very inefficient. We need a better filesystem ... so I started working on FoxFS and got it to a state worth specifying after just 3 days. It is not yet completed and please don't trust it to store your valuable data, but please test it and report errors so I can fix them.

1.2 Design

FoxFS has a layered design: below FoxFS is the block-access layer provided by the device-driver.

On top of this is the blockchain-layer, this is the first layer in FoxFS and implements data-storage inside blockchains (not the Bitcoin-ones, but just a chain of blocks e.g. first load block 5, then 7 and then 6).

On the foundation of the blockchain-layer, FoxFS builds the filesystem-layer, this is where the actual filesystem is implemented. In the filesystem-layer we have directories, files, file-attributes, ... everything you would expect from a filesystem.

2 Datastructures in FoxFS

2.1 superblock

The superblock contains indispensable filesystem management information. It starts with a bit of magic(numbers) to make identifying the filesystem type easy, followed by a version number(currently 1). Incompatible changes to filesystem structures should be accompanied by an incremented version number. If a filesystem driver encounters a version number it doesn't know, it must not execute any write operations and emit a warning. It may however read the contents of the filesystem (possibly wrongly) or it may emit an error message and not read the filesystem at all.

The superblock also contains the blocksize of the filesystem. A driver should obey the blocksize specified in the superblock unless the user overrides it.

The last components of the superblock are "pointers"(block ids and offsets) to special blockchains.

The superblock must always be located in block 0. The presence and behavior of additional copies is unspecified.

2.2 inode

An inode is a simple structure, which points to a continuous area of data-blocks and the following inode. It is used to build the blockchains, where we put all the data, and the free space map which keeps track of unused blocks. The format is shown in table 2.

Table 1: Format of the superblock

Field	Datatype	Description
magic	char[4]	magic number 'FFSS' (FoxFS Superblock)
version	unsigned 8bit integer	the filesystem version
reserved	24 bit	reserved for future expansion
block_size	unsigned 64bit integer	
root_dir_block_idx	unsigned 64bit integer	
root_dir_offset	unsigned 32bit integer	
free_space_chain_block_idx	unsigned 64bit integer	
free_space_chain_offset	unsigned 32bit integer	
fs_metadata_block_idx	unsigned 64bit integer	
fs_metadata_offset	unsigned 32bit integer	

Table 2: Format of an inode

Field	Datatype	Description
block_idx	unsigned 64bit integer	The block containing the data for this inode
length	unsigned 64bit integer	the length of this block area in blocks
next_inode_block_idx	unsigned 64bit integer	The block containing the next inode
next_inode_offset	unsigned 32bit integer	The offset of the inode in the block, given in bytes
magic	4 bytes	Magic number should be set to 'NODE' on write and checked on read

The data associated with the inode can be found in the blocks referenced by `block_idx` and `length`. The next inode can be found with the `next_inode_*`-fields, if there is no following inode, these are set to 0.

To get all data in a blockchain, read the blocks referenced by the first inode, get the next inode, load the blocks referenced in that inode, get the next inode, ... until there is no following inode.

Inodes are identified by the block where they are stored and the offset from the beginning of the block given as `inode-index`.

2.3 free space chain

A special blockchain which does not contain data but instead only references free blocks. Like ordinary blockchains, it is composed of Inodes and can be used to find an available block to store new data.

A block can be unavailable for many reasons, for example damage or it is just already used. In either way, it is always just marked as unavailable, it's not our department to track disk-health.

2.4 Directories

This is the core-function of every filesystem: mapping filenames and metadata to actual filecontent.

Directories in FoxFS are stored in a blockchain. Actually, there isn't any difference between directories and files. Directories are just files with a fileformat known by the filesystem-driver.

The format of directories is specified in table 3. All entries are saved without holes in the structure, making the files less fragmented and less space-consuming. This way, we only need to save the number of entries and not their position or a valid flag in every entry. If a driver deletes an entry, it has to fill the hole with other entries. The order of the entries is guaranteed not to be preserved.

As you can see in table 4, a directory-entry only contains references to inodes. There is one inode for the data-blockchain, which contains the actual file content and another inode for metadata-blockchain.

With this design, we can store an arbitrary amount of metadata in the filesystem, much more than the normal filename, size and different timestamps - we could store changelogs or diffs there for example.

Table 3: Structure of a directory

Field	Datatype	Description
num_entries	unsigned 64bit integer	The number of entries in this directory
children_size	unsigned 64bit integer	Last known size of all children
entries	array of structure defined in table 4	The actual entries in this directory

2.5 File-metadata

This is the structure stored in the meta-blockchain and referenced by every file.

3 The layers of FoxFS

3.1 Block-layer

This is actually not a layer from FoxFS, but the layer below it. FoxFS needs just three functions in this layer: `storeBlock`, `loadBlock` and `getNumBlocks` - everything else is implemented on these. FoxFS also supports the `trimBlock` function, enabling the block-layer to know which blocks are free.

3.2 Blockchain-layer

This is the first layer in FoxFS. It implements chains of blocks to store data in. Everything in FoxFS is stored in such a blockchain.

Blockchains are identified by the inode pointing to their first block.

Table 4: Structure of an fileentry

Field	Datatype	Description
content_inode_block_idx	unsigned 64bit integer	Block containing the inode for the content-blockchain
content_inode_offset	unsigned 32bit integer	Offset of the inode in the block given in content_inode_block_idx
meta_inode_block_idx	unsigned 64bit integer	Block containing the inode for the metadata-blockchain
meta_inode_offset	unsigned 32bit integer	Offset of the inode in the block given in meta_inode_block_idx
filename	unsigned 8bit integer[30]	Name of the file. If the name is longer, it is stored in the metadata and this field is filled with zeros. The name is stored as UTF-8 and without trailing zero.
size	unsigned 64bit integer	Size of the file given in bytes
attributes	unsigned 16bit integer	some attributes: <ul style="list-style-type: none"> • 0x01 Entry is a directory • 0x02 Entry is a file
reserved	unsigned 8bit integer	reserved for future use. Set to 0 on write and keep value on update.

Table 5: Metadata structure

Field	Datatype	Description
magic	unsigned 32bit integer	'M3TA'
num_entries	unsigned 64bit integer	Number of metadata-entries